

In [33]:

```
1
2 %matplotlib inline
3
4 import numpy as np
5 import pandas as pd
6 import scipy.stats as stats
7 import matplotlib.pyplot as plt
8 import math
9 import random
10 from matplotlib import collections as matcoll
11 import seaborn as sns
12 from scipy.optimize import curve_fit
13 from sklearn.linear_model import LinearRegression
14 from pylab import *
15 from kapteyn import kmpfit
16 import pylab
17
18 import statsmodels.api as sm
19 from statsmodels.stats.outliers_influence import summary_table
```

In [3]:


```
1 #Read Data into dataframe
2 gaugeData = pd.read_table('gauge-1wb1wa6-2gpel41.txt',delim_whitespace=True)
```

```
In [4]:
```

```
1 plt.scatter(gaugeData['density'], gaugeData['gain'])
```

**Out[4]:**

<matplotlib.collections.PathCollection at 0x1a14a4e7b8>



```
In [5]:
```


```
1 #Log Transformation Gain Data
2 gaugeData2 = gaugeData.copy(deep = True)
3 gaugeData2['gain'] = gaugeData2['gain'].apply(lambda x: np.log10(x))
```

In [6]:

```
1 plt.scatter(gaugeData2['density'], gaugeData2['gain'])
```


Out[6]:

<matplotlib.collections.PathCollection at 0x1a14ba6278>




In [7]:

```
1 #Data plotted with least squares line
2 def line(x, m, b):
3     return m*x + b
4 M,B = curve_fit(line, gaugeData['density'], gaugeData['gain'])[0]
5 def abline(slope, intercept):
6     axes = plt.gca()
7     x_vals = np.array(axes.get_xlim())
8     y_vals = intercept + slope * x_vals
9     plt.plot(x_vals, y_vals, '--')
10 plt.scatter(gaugeData['density'],gaugeData['gain'])
11 plt.ylim(0,400)
12 abline(M,B)
```



In [8]:

```
1 #Data plotted with least squares line
2 def line(x, m, b):
3     return m*x + b
4 M2,B2 = curve_fit(line, gaugeData2['density'], gaugeData2['gain'])[0]
5 def abline(slope, intercept):
6     axes = plt.gca()
7     x_vals = np.array(axes.get_xlim())
8     y_vals = intercept + slope * x_vals
9     plt.plot(x_vals, y_vals, '--')
10 plt.scatter(gaugeData2['density'],gaugeData2['gain'])
11 plt.ylim(0,2.8)
12 abline(M2,B2)
```




In [9]:

```
1 #log transformed residuals
2 sns.residplot(gaugeData2['density'],gaugeData2['gain'])
```

Out[9]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1a151e9710>




In [10]:

```
1 #Original Data Residuals
2 sns.residplot(gaugeData['density'],gaugeData['gain'])
```

Out[10]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1a152db278>



In [11]:


```
1 np.correlate(gaugeData['density'], gaugeData['gain'])
```

Out[11]:

```
array([ 1824.5634])
```

In [35]:

```
1 #Different Model:
2 model2 = LinearRegression(fit_intercept=True)
3
4 model2.fit(gaugeData2['density'][:, np.newaxis], gaugeData2['gain'])
5
6 xfit = np.linspace(0, 1, 10)
7 yfit = model2.predict(xfit[:, np.newaxis])
8
9 plt.scatter(gaugeData2['density'], gaugeData2['gain'])
10 plt.xlim(0, 0.7)
11 plt.plot(xfit, yfit);
```




In [28]:

```
1 x = gaugeData['density']
2 y = gaugeData['gain']
3
4 X = sm.add_constant(x)
5 res = sm.OLS(y, X).fit()
6
```

```

7 st, data, ss2 = summary_table(res, alpha=0.05)
8 fitvalues = data[:,2]
9 predict_mean_se = data[:,3]
10 predict_mean_l, predict_mean_u = data[:,4:6].T
11 predict_l, predict_u = data[:,6:8].T
12
13 fig, ax = plt.subplots(figsize=(8,6))
14 ax.plot(x, y, 'o', label="data")
15 ax.plot(X, fitvalues, 'r-', label='OLS')
16 ax.plot(X, predict_l, 'b--')
17 ax.plot(X, predict_u, 'b--')
18 ax.plot(X, predict_mean_l, 'g--')
19 ax.plot(X, predict_mean_u, 'g--')
20 plt.xlim(0,0.7)
21 plt.show()

```



In [148]:

```

1 x = gaugeData2['density']
2 y = gaugeData2['gain']
3
4 X = sm.add_constant(x)


```



```

5 res = sm.OLS(y, X).fit()
6
7 st, data, ss2 = summary_table(res, alpha=0.05)
8 fitvalues = data[:,2]
9 predict_mean_se = data[:,3]
10 predict_mean_l, predict_mean_u = data[:,4:6].T
11 predict_l, predict_u = data[:,6:8].T
12
13 fig, ax = plt.subplots(figsize=(8,6))
14 ax.plot(x, y, 'o')
15 ax.plot(X, fitvalues, 'r-')
16 ax.plot(X, predict_l, 'b--')
17 ax.plot(X, predict_u, 'b--')
18 ax.plot(X, predict_mean_l, 'g--')
19 ax.plot(X, predict_mean_u, 'g--')
20 plt.xlim(0,0.7)
21 plt.show()

```



In [161]:

```
1  # Cross Validation with
2  std_error = 10**.0138181
3  def cross_validate(density):
4      sumGain = (gaugeData2.gain[gaugeData2.density == density].sum())
5      lenGain = len(gaugeData2.gain[gaugeData2.density == density])
6      meanGain = sumGain/lenGain
7      nDF = gaugeData2[gaugeData2.density != density]
8      M3,B3 = curve_fit(line, nDF['density'], nDF['gain'])[0]
9      yr = M3*density + B3
10     xr = (np.log(38.6) - B3)/M3
11     return xr
```